

1. Benefits of the Relational Model

- Standard data model and query language
- ACID compliance (Atomicity, Consistency, Isolation, Durability)
- Works well with highly structured data
- Can handle large volumes of data
- Well-supported with tools and experience

2. Relational Database Performance Enhancements

- Indexing
- Storage control (column vs. row orientation)
- Query optimization
- Caching & prefetching
- Materialized views
- Precompiled stored procedures
- Data replication & partitioning

3. Transaction Processing

- **Transaction:** A sequence of CRUD operations treated as one logical unit
 - Entire sequence **commits** or **rolls back**
- Ensures:
 - Data integrity
 - Error recovery
 - Concurrency control
 - Reliable data storage
 - Simplified error handling

4. ACID Properties

- **Atomicity:** Transaction fully completes or doesn't execute at all
- **Consistency:** Transitions DB from one valid state to another
- **Isolation:** Transactions don't interfere with each other
 - Issues:
 - Dirty Reads: Read uncommitted data
 - Non-repeatable Reads: Same query returns different results
 - Phantom Reads: Row set changes during transaction
- **Durability:** Once committed, changes are permanent—even after failure

5. SQL Transaction Example: Transfer Money

- Transfers money between two accounts with rollback if funds are insufficient
- Uses:
 - START TRANSACTION
 - UPDATE accounts
 - Conditional check
 - ROLLBACK or COMMIT
 - Inserts into a transactions table

6. Why Move Beyond Relational?

- Schema evolution over time
- Not all applications require full ACID
- Expensive joins
- Increasing semi-structured/unstructured data (e.g., JSON, XML)
- Difficult horizontal scaling
- Need for low latency / real-time systems

7. Scalability: Vertical vs Horizontal

- **Scale Up (Vertical):**
 - Bigger servers, simpler
 - Eventually hits cost/performance limits
- **Scale Out (Horizontal):**
 - Distributed computing models
 - Modern tools make this easier

8. Distributed Systems

- “A collection of independent computers that appear as one”
(*Andrew Tennenbaum*)
- Key properties:
 - Concurrent operation
 - Independent failures
 - No global clock

9. Distributed Data Stores

- Data stored on multiple nodes (replicated or partitioned)
- Examples:
 - MySQL/PostgreSQL (support replication/sharding)
 - CockroachDB
 - Many NoSQL databases
- Systems must be **Partition Tolerant**
 - Must operate despite network/system failures

10. CAP Theorem

- Cannot achieve all three simultaneously:
 1. **Consistency**: Latest write always read
 2. **Availability**: Every request gets a response
 3. **Partition Tolerance**: Operates despite network issues

11. CAP Theorem Trade-offs

- **Consistency + Availability**: Not partition-tolerant
- **Consistency + Partition Tolerance**: May sacrifice availability
- **Availability + Partition Tolerance**: May sacrifice consistency

Note: CAP's "Consistency" \neq ACID's "Consistency"

12. CAP Theorem Real-World Takeaway

- You must trade off between:
 - Consistency
 - Availability
 - Partition Tolerance
- You can't have all three under fault conditions