

1. ACID and Concurrency

Pessimistic Concurrency (Traditional ACID)

- Focuses on **data safety**
- Assumes things will go wrong
- Prevents conflicts via **locks** (read/write)
 - Analogy: borrowing a library book – exclusive access

Optimistic Concurrency

- Assumes **conflicts are rare**
- Doesn't lock on read/write
- Uses **timestamps and version numbers** to detect conflict at commit time
- Works well in **read-heavy/low-conflict** environments (e.g., analytics)
- Locking is better for **high-conflict** systems

2. What is NoSQL?

- Originally meant “No SQL,” now means “Not Only SQL”
- Often refers to **non-relational databases**
- Created to handle **semi/unstructured web data**

3. CAP Theorem (Review)

- A distributed system **can't guarantee all three**:
 - **Consistency**: Every read reflects the latest write
 - **Availability**: System continues to respond to all requests
 - **Partition Tolerance**: System operates despite network failures
- Trade-offs:
 - C + A: Can't tolerate partitioning
 - C + P: Might drop availability
 - A + P: May serve stale data

4. BASE – The ACID Alternative for Distributed Systems

- **Basically Available:** System is usually available
- **Soft State:** State can change over time without input
- **Eventual Consistency:** System will become consistent over time

5. Key-Value (KV) Databases

- Simple **key = value** data model
- Prioritizes:
 - **Simplicity:** CRUD ops and API-friendly
 - **Speed:** Often in-memory; $O(1)$ access using hash tables
 - **Scalability:** Easy horizontal scaling; embraces eventual consistency

6. Use Cases for Key-Value Stores

- **Data Science:** EDA results, experiment tracking
- **ML:** Feature stores, real-time model metrics
- **Web Dev:**
 - Session storage
 - User preferences
 - Shopping carts
 - Caching layer

7. Redis – A Popular Key-Value Store

- **Open-source, in-memory database**
- Also supports:
 - Graphs, Full Text Search, Time Series, Geospatial
- **Durability** via:
 - Snapshots
 - Append-only file
- Over 100k+ ops/sec
- Lookup only by key (no secondary indexes)

8. Redis Data Types

- **String**: basic type, maps string → string
- **Hash**: object-like structures (field-value pairs)
- **List**: linked list of strings; supports stacks & queues
- **Set**: unordered, unique strings; supports set operations
- **JSON**: full support, uses JSONPath for querying

9. Redis Setup

- **Docker**:
 - Search + run Redis image
 - Optional: expose port 6379 (default Redis port)
- **DataGrip**:
 - File > New > Data Source > Redis
 - Set port to 6379 and test connection

10. Redis Basics – String Type

- Use for caching, config settings, token mgmt, counters
- Example Commands:
 - SET key value, GET key, EXISTS key, DEL key, KEYS pattern
 - INCR key, DECRBY key amt
 - SETNX key value: only sets if key doesn't exist

11. Redis Hash Type

- Use to represent objects (e.g., a bike with model, brand, price)
- Example Commands:
 - HSET bike:1 model Demios
 - HGETALL bike:1, HINCRBY bike:1 price 100

12. Redis List Type

- Great for:
 - Stacks, queues, logging, task queues, chat/message history
- Key Commands:
 - LPUSH, RPUSH, LPOP, RPOP
 - LRange key start stop, LLEN key

13. Redis Set Type

- Use for:
 - Tracking unique values (IP addresses, user IDs)
 - Access control, friends lists, groups
- Supports full **set ops**:
 - SADD, SISMEMBER, SCARD
 - SINTER, SDIFF, SREM, SRANDMEMBER